

OM CHOKSI

🎯 Learning Objectives

- Document Classification
- Text Pre-processing
- Feature extraction
- Vocabulary creation

📋 Vocabulary & Feature Extraction

Given a document, you can represent it as a vector of dimension V , where V corresponds to your vocabulary size. As V gets larger, the vector becomes more sparse. Furthermore, we end up having many more features and end up training lot of parameters. This could result in larger training time, and large prediction time.

🛠️ Preprocessing

When preprocessing, you have to perform the following:

1. Eliminate handles and URLs
2. Tokenize the string into words
3. Remove stop words like "and, is, a, on, etc."
4. Stemming - or convert every word to its stem. Like dancer, dancing, danced, becomes 'danc'.
5. Convert all your words to lower case.

📌 Cell 1: Imports & Setup

```
In [1]: import warnings  
warnings.filterwarnings("ignore")  
  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
  
from sklearn.datasets import fetch_20newsgroups  
from sklearn.decomposition import TruncatedSVD
```

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import accuracy_score, classification_report, confusion_m
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import Normalizer
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC

import tensorflow as tf
from tensorflow.keras import layers
import tensorflow_datasets as tfds

```

```

2025-11-17 19:04:48.094720: E external/local_xla/xla/stream_executor/cuda/cud
a_fft.cc:477] Unable to register cuFFT factory: Attempting to register factory
for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
E0000 00:00:1763406288.621020      48 cuda_dnn.cc:8310] Unable to register cuDN
N factory: Attempting to register factory for plugin cuDNN when one has already
been registered
E0000 00:00:1763406288.734810      48 cuda_blas.cc:1418] Unable to register cuB
LAS factory: Attempting to register factory for plugin cuBLAS when one has alre
ady been registered

-----
AttributeError                                     Traceback (most recent call last)
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'

-----
AttributeError                                     Traceback (most recent call last)
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'

-----
AttributeError                                     Traceback (most recent call last)
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'

-----
AttributeError                                     Traceback (most recent call last)
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'

-----
AttributeError                                     Traceback (most recent call last)
AttributeError: 'MessageFactory' object has no attribute 'GetPrototype'

```

Cell 2: Load Data

```

In [2]: categories = ['alt.atheism', 'talk.religion.misc', 'comp.graphics', 'sci.space']

data_train = fetch_20newsgroups(subset='train', categories=categories, shuffle=True)
data_test = fetch_20newsgroups(subset='test', categories=categories, shuffle=True)

print("Train samples:", len(data_train.data))
print("Test samples:", len(data_test.data))

y_train, y_test = data_train.target, data_test.target
target_names = data_train.target_names

```


```
Train samples: 2034  
Test samples: 1353
```

Cell 3: Feature Extraction

```
In [3]: vectorizer = TfidfVectorizer(max_df=0.5, min_df=2, stop_words='english')  
X_train = vectorizer.fit_transform(data_train.data)  
X_test = vectorizer.transform(data_test.data)  
  
# Optional dimensionality reduction  
n_components = 100  
svd = TruncatedSVD(n_components)  
normalizer = Normalizer(copy=False)  
lsa = make_pipeline(svd, normalizer)  
  
X_train = lsa.fit_transform(X_train)  
X_test = lsa.transform(X_test)
```

Cell 4: Logistic Regression

```
In [4]: lr_clf = LogisticRegression(max_iter=1000)  
lr_clf.fit(X_train, y_train)  
  
lr_train_pred = lr_clf.predict(X_train)  
lr_test_pred = lr_clf.predict(X_test)  
  
print("Logistic Regression Train Accuracy:", accuracy_score(y_train, lr_train_<pre>print("Logistic Regression Test Accuracy:", accuracy_score(y_test, lr_test_<pre>ConfusionMatrixDisplay(confusion_matrix(y_test, lr_test_pred), display_labels=<pre>plt.title("Logistic Regression Confusion Matrix")  
plt.show()  
  
print(pd.DataFrame(classification_report(y_test, lr_test_pred, output_dict=True))<pre>Logistic Regression Train Accuracy: 0.95968534906588  
Logistic Regression Test Accuracy: 0.8787878787878788
```



	precision	recall	f1-score	support
0	0.817035	0.811912	0.814465	319.000000
1	0.939086	0.951157	0.945083	389.000000
2	0.928218	0.951777	0.939850	394.000000
3	0.777311	0.737052	0.756646	251.000000
accuracy	0.878788	0.878788	0.878788	0.878788
macro avg	0.865412	0.862974	0.864011	1353.000000
weighted avg	0.877133	0.878788	0.877805	1353.000000

Cell 5: Naive Bayes

```
In [5]: nb_clf = GaussianNB()
nb_clf.fit(X_train, y_train)


nb_train_pred = nb_clf.predict(X_train)
nb_test_pred = nb_clf.predict(X_test)

print("Naive Bayes Train Accuracy:", accuracy_score(y_train, nb_train_pred))
print("Naive Bayes Test Accuracy:", accuracy_score(y_test, nb_test_pred))

ConfusionMatrixDisplay(confusion_matrix(y_test, nb_test_pred), display_labels=
plt.title("Naive Bayes Confusion Matrix")
plt.show()

print(pd.DataFrame(classification_report(y_test, nb_test_pred, output_dict=True))
```

```
Naive Bayes Train Accuracy: 0.8593903638151426  
Naive Bayes Test Accuracy: 0.7871396895787139
```




Cell 6: Support Vector Machine

```
In [6]: svm_clf = SVC()  
svm_clf.fit(X_train, y_train)  
  
svm_train_pred = svm_clf.predict(X_train)  
svm_test_pred = svm_clf.predict(X_test)  
  
print("SVM Train Accuracy:", accuracy_score(y_train, svm_train_pred))  
print("SVM Test Accuracy:", accuracy_score(y_test, svm_test_pred))  
  
ConfusionMatrixDisplay(confusion_matrix(y_test, svm_test_pred), display_labels  
plt.title("SVM Confusion Matrix")  
plt.show()
```

```

print(pd.DataFrame(classification_report(y_test, svm_test_pred, output_dict=True))
SVM Train Accuracy: 0.9837758112094396
SVM Test Accuracy: 0.8906134515890614

```



	precision	recall	f1-score	support
0	0.847403	0.818182	0.832536	319.000000
1	0.933002	0.966581	0.949495	389.000000
2	0.944162	0.944162	0.944162	394.000000
3	0.790323	0.780876	0.785571	251.000000
accuracy	0.890613	0.890613	0.890613	0.890613
macro avg	0.878723	0.877450	0.877941	1353.000000
weighted avg	0.889601	0.890613	0.889956	1353.000000

Cell 7: Deep Learning CNN

```

In [7]: tokenizer = tfds.deprecated.text.SubwordTextEncoder.build_from_corpus(
    data_train.data, target_vocab_size=2**15
)

train_inputs = [tokenizer.encode(text) for text in data_train.data]
test_inputs = [tokenizer.encode(text) for text in data_test.data]

MAX_LEN = max(len(x) for x in train_inputs)

train_inputs = tf.keras.preprocessing.sequence.pad_sequences(train_inputs, val

```

```

test_inputs = tf.keras.preprocessing.sequence.pad_sequences(test_inputs, value=0)

class DCNN(tf.keras.Model):
    def __init__(self, vocab_size, emb_dim=128, nb_filters=50, FFN_units=512,
                 super(DCNN, self).__init__()
                 self.embedding = layers.Embedding(vocab_size, emb_dim)
                 self.bigram = layers.Conv1D(filters=nb_filters, kernel_size=2, activation="relu")
                 self.trigram = layers.Conv1D(filters=nb_filters, kernel_size=3, activation="relu")
                 self.fourgram = layers.Conv1D(filters=nb_filters, kernel_size=4, activation="relu")
                 self.pool = layers.GlobalMaxPool1D()
                 self.dense_1 = layers.Dense(units=FFN_units, activation="relu")
                 self.dropout = layers.Dropout(rate=dropout_rate)
                 self.last_dense = layers.Dense(units=nb_classes, activation="softmax")

    def call(self, inputs, training=False):
        x = self.embedding(inputs)
        x_1 = self.pool(self.bigram(x))
        x_2 = self.pool(self.trigram(x))
        x_3 = self.pool(self.fourgram(x))
        merged = tf.concat([x_1, x_2, x_3], axis=-1)
        merged = self.dense_1(merged)
        merged = self.dropout(merged, training=training)
        return self.last_dense(merged)

VOCAB_SIZE = tokenizer.vocab_size
Dcnn = DCNN(vocab_size=VOCAB_SIZE, emb_dim=128, nb_filters=100, FFN_units=256,
             super(DCNN, self).__init__()
             self.embedding = layers.Embedding(vocab_size, emb_dim)
             self.bigram = layers.Conv1D(filters=nb_filters, kernel_size=2, activation="relu")
             self.trigram = layers.Conv1D(filters=nb_filters, kernel_size=3, activation="relu")
             self.fourgram = layers.Conv1D(filters=nb_filters, kernel_size=4, activation="relu")
             self.pool = layers.GlobalMaxPool1D()
             self.dense_1 = layers.Dense(units=FFN_units, activation="relu")
             self.dropout = layers.Dropout(rate=dropout_rate)
             self.last_dense = layers.Dense(units=nb_classes, activation="softmax")

history = Dcnn.fit(train_inputs, y_train, batch_size=16, epochs=10, validation_data=(val_inputs, val_labels))

```

```

I0000 00:00:1763406543.139391      48 gpu_device.cc:2022] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 13942 MB memory: -> device: 0, name: Tesla T4, pci bus id: 0000:00:04.0, compute capability: 7.5
I0000 00:00:1763406543.140068      48 gpu_device.cc:2022] Created device /job:localhost/replica:0/task:0/device:GPU:1 with 13942 MB memory: -> device: 1, name: Tesla T4, pci bus id: 0000:00:05.0, compute capability: 7.5
Epoch 1/10

```

```

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
I0000 00:00:1763406547.051772      194 service.cc:148] XLA service 0x7c071000954
0 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:
I0000 00:00:1763406547.053060      194 service.cc:156]     StreamExecutor device
(0): Tesla T4, Compute Capability 7.5
I0000 00:00:1763406547.053076      194 service.cc:156]     StreamExecutor device
(1): Tesla T4, Compute Capability 7.5
I0000 00:00:1763406547.446870      194 cuda_dnn.cc:529] Loaded cuDNN version 903
00
    1/128 ━━━━━━━━━━ 20:51 10s/step - loss: 1.4164 - sparse_categorical_accuracy: 0.1250
I0000 00:00:1763406553.156719      194 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.


```

```
128/128 32s 173ms/step - loss: 1.3169 - sparse_categorical_accuracy: 0.3599 - val_loss: 0.7359 - val_sparse_categorical_accuracy: 0.6911
Epoch 2/10
128/128 17s 134ms/step - loss: 0.3814 - sparse_categorical_accuracy: 0.8751 - val_loss: 0.4556 - val_sparse_categorical_accuracy: 0.8551
Epoch 3/10
128/128 17s 136ms/step - loss: 0.0463 - sparse_categorical_accuracy: 0.9949 - val_loss: 0.4047 - val_sparse_categorical_accuracy: 0.8736
Epoch 4/10
128/128 18s 139ms/step - loss: 0.0108 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.4335 - val_sparse_categorical_accuracy: 0.8729
Epoch 5/10
128/128 19s 145ms/step - loss: 0.0050 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.4500 - val_sparse_categorical_accuracy: 0.8758
Epoch 6/10
128/128 19s 147ms/step - loss: 0.0026 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.4627 - val_sparse_categorical_accuracy: 0.8788
Epoch 7/10
128/128 18s 143ms/step - loss: 0.0015 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.4848 - val_sparse_categorical_accuracy: 0.8758
Epoch 8/10
128/128 18s 142ms/step - loss: 0.0015 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.4959 - val_sparse_categorical_accuracy: 0.8773
Epoch 9/10
128/128 18s 143ms/step - loss: 7.1925e-04 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.5023 - val_sparse_categorical_accuracy: 0.8758
Epoch 10/10
128/128 18s 144ms/step - loss: 0.0012 - sparse_categorical_accuracy: 1.0000 - val_loss: 0.5312 - val_sparse_categorical_accuracy: 0.8780
```

Cell 8: Plot Training Curves

```
In [8]: plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.legend(); plt.title("Loss Evolution")

plt.subplot(1,2,2)
plt.plot(history.history['sparse_categorical_accuracy'], label='Train Acc')
plt.plot(history.history['val_sparse_categorical_accuracy'], label='Val Acc')
plt.legend(); plt.title("Accuracy Evolution")
plt.show()
```



Cell 9: Final Evaluation


```
In [9]: results = Dcnn.evaluate(test_inputs, y_test, batch_size=16)
print("Final Test Loss:", results[0])
print("Final Test Accuracy:", results[1])

85/85 ██████████ 2s 24ms/step - loss: 0.4986 - sparse_categorical_acc
uracy: 0.8763
Final Test Loss: 0.5311814546585083
Final Test Accuracy: 0.8780487775802612
```

Cell 10: Comparison Bar Chart

```
In [10]: results_dict = {
    "Logistic Regression": accuracy_score(y_test, lr_test_pred),
    "Naive Bayes": accuracy_score(y_test, nb_test_pred),
    "SVM": accuracy_score(y_test, svm_test_pred),
    "CNN": results[1]
}

plt.figure(figsize=(8,6))
plt.bar(results_dict.keys(), results_dict.values(), color=['blue','green','orange'])
plt.ylabel("Test Accuracy")
plt.title("Model Comparison on 20 Newsgroups Subset")
plt.ylim(0,1)
for i, v in enumerate(results_dict.values()):
    plt.text(i, v+0.01, f"{v:.2f}", ha='center')
plt.show()
```



In []: